

Command/Concept	Usage/Syntax	Description	Example
qsub	qsub <script_to_run>	Submit a "batch" job to the SGE queuing system ("the hard way"). See the CGRB's SGE_Batch (below) for an easier submission method. SGE stdout (*.o<job id>) and stderr (*.e<job id>) files will be placed in your \$HOME directory by default.	# For numerous options qsub -help # Simple script cat sgetest.sh #!/bin/tcsh echo "SGE Test!" # Simple submission qsub sgetest.bash
qstat	qstat	Status of "running" jobs in SGE. "state" will show: r Job is running on a compute node qw Waiting in queue to be matched to a node eqw Error, delete job and try resubmitting Completed jobs will <u>no longer appear</u> in qstat	# Show your running jobs qstat # Show <u>all</u> jobs on system qstat -u '*'
qdel	qdel <job number> qdel -u <username>	Delete a submitted job. You can only delete your jobs, not another user's jobs.	# Deletes a job you own qdel <job number> # Deletes <u>all</u> of your jobs qdel -u <username>
qhost SGE_Avail	qhost qhost -q SGE_Avail	Show resources (machines and queues) available in SGE. The CGRB's SGE_Avail wrapper script shows resources and queues available <u>specifically to you</u> .	# Show resources for <u>all</u> # available machines in SGE qhost # All resources by queue qhost -q # Resources available to <u>you</u> SGE_Avail

qrsh	qrsh -l h=" <hostname> " qrsh -pe thread <int> -l h=" <hostname> "	Interactively login to a machine reserving 1 CPU (thread) by default from SGE. Note: This is turned <u>off</u> on the majority of CGRB machines. If your lab or department owns a machine you may be granted access to use qrsh to login interactively to it.	# Interactively login to # 'server100' reserving 1 CPU qrsh -l h="server100" # Interactively login to # 'server100' reserving 4 # CPUs from SGE qrsh -pe thread 4 -l h="server100"
Queue types	SGE_Avail or qhost -q	A machine can have 3 queue properties under "ARCH" or "QTYPE": B Allow for "batch" submission via qsub or SGE_Batch I Allow for interactive use with qrsh P Allow for multiple-CPU jobs (almost all of them)	SGE_Avail or qhost -q
SGE_Batch	Interactive: SGE_Batch Command-line: SGE_Batch -P <reserve # CPUs> -f <free RAM> -m <max RAM> -F <max file size> -q <queue to use> -c <command to run> -r <run output dir>	CGRB's qsub wrapper script ("the easy way" to submit jobs to SGE). Running the command by itself loads an interactive menu; options c (command) and r (run output directory) are <u>required</u> . h to show setup and s submits job. * .sh (script file created by SGE_Batch), SGE stdout and stderr files are placed in the run output directory you specify and <u>not</u> your \$HOME	# See usage/examples SGE_Batch -help # Submit sgetest.sh # job requesting resources: # 4 CPU slots, 5 GB of RAM # Kill if uses > 5 GB of RAM # Kill if it creates a single file # greater than 100 GB # Put logs in \$PWD/mytest SGE_Batch -c sgetest.sh -P 4 -f 5G -m 5G -F 100G -r mytest

Array jobs	<code>SGE_Batch -t <range></code>	SGE_Batch can submit an "array of jobs" from one script. You specify the range of numbers via <code>-t</code> (or <code>t</code> interactively), e.g., <code>-t 1-20</code> or if by twos, <code>2-20:2</code> The total number of jobs launched by SGE_Batch is equal to the range of numbers you specified, e.g., 20 Each SGE job of the "array" is assigned one of the numbers and accessible in your script via <code>\$SGE_TASK_ID</code> All jobs will write their SGE <code>stderr</code> and <code>stdout</code> files to the same run output directory	# Simple script echoing its # assigned ID <code>cat sgearray.sh</code> <code>#!/bin/tcsh</code> <code>echo "My assigned ID number: \$SGE_TASK_ID"</code> <code>sleep 60</code> # Submit 10 array jobs # Using 1 CPU and 1 GB of # RAM each <code>SGE_Batch</code> <code>-c sgearray.sh</code> <code>-t 1-10</code> <code>-P 1 -f 1G -m 1G</code> <code>-r arraytest</code> # Repeatedly run until you # see all 10 jobs queued up <code>Qstat</code>
SGE_Plot	<code>SGE_Plotdir</code> <code><run output dir></code>	Displays RAM usage and execution time. If the run output directory contains the results of an "Array Jobs" run it will produce ASCII plots of this information.	# Plot the previous array jobs # test (producing 10 jobs) <code>SGE_Plotdir arraytest</code> <code> less</code>
"10/100/1000" rule	<code>SGE_Batch</code> and <code>SGE_Plotdir</code>	Run tests on <u>small</u> subsets of data to extrapolate usage requirements! Bioinformatic algorithms behave differently and <u>may scale</u> in different ways, e.g., linearly, logarithmically, etc. Test, test, test!	Run SGE_Batch using "Array Jobs" with different sized data sets, e.g., 1/1000 th , 1/100 th , 1/10 th to compare memory usage and CPU run times, e.g., via SGE_Plotdir